

PRODUCT OVERVIEW

RAG-OS: the knowledge assistant that knows who is asking

Answers drawn only from your documents, only from the ones each person is allowed to see, with a page-level citation for every claim, and an honest "I don't know" when the documents don't say.

Version 1.0 **Date** 2 October 2026 **Platform** Microsoft Azure · Microsoft Foundry

Grounded

Every answer cites its source passages; uncited answers are refused.

Permission-aware

Department, region, clearance and role are enforced inside the search.

Yours

Runs in your Azure subscription, with models you choose.

Accountable

Token cost and a correlation ID on every answer.

How to read this document:

AVAILABLE

built and running today

PARTIAL

works, with known limits stated

ROADMAP

planned, not yet built

Contents

1. Executive summary	3
2. Introduction	4
3. Product overview	5
3.1 Data ingestion path	7
3.2 User query path	9
3.3 Answer quality scenarios	11
4. Models	12
5. Secure data querying	13
6. Data compliance	16
7. Observability and operations	17
8. How RAG-OS is different	18
9. Similar products in market	20
10. Potential use cases	22
11. Future scope and potential	24
12. Appendix	25

1 Executive summary

Every organisation already holds the answers its people need. They sit in policies, price lists, SOPs, contracts and handbooks spread across HR, Sales, IT, Legal and Finance, and finding the right one, in its current version, is slow and error-prone.

Generic AI assistants made that search feel instant, but they introduced two risks that enterprises cannot accept:

- **They leak.** An assistant that can read everything will eventually quote a confidential salary band to an intern or a restricted contract clause to a partner.
- **They guess.** When the documents do not cover a question, a language model still produces a fluent, plausible answer, and nobody can tell which parts are real.

RAG-OS removes both risks by design. Each user's identity attributes (department, region, clearance level and role) are compiled into a filter that runs *inside* the search engine, so documents a person may not see are never retrieved, never sent to the model and never paraphrased. Answers are generated only from retrieved passages, carry numbered citations down to the page, and are refused outright when the evidence is missing.

The business case

Faster answers, fewer escalations

Staff self-serve policy, pricing and process questions that today land on HR, Sales Ops, IT or Legal, with a citation they can open and verify.

Lower risk of disclosure

Access is enforced before generation, not filtered afterwards. The same assistant can safely serve an entire workforce, and even external portal users.

Predictable, visible cost

Embeddings run on self-hosted models with no per-token charge, and every answer reports its own token usage, so cost per question is measurable from day one.

A platform, not a chatbot

Once the knowledge base is indexed and classified, the same governed retrieval can power questionnaires, form filling and vertical workflows (see section 10).

Why now

Organisations have moved from AI pilots to production rollouts, and the questions buyers ask have changed accordingly: *who can see what, where does the data go, can we prove where an answer came from and what does each answer cost*. RAG-OS is built around exactly those questions, on the Azure and Microsoft Foundry estate many enterprises have already approved.

2 Introduction

RAG-OS is an enterprise-grade knowledge-base assistant built on **retrieval-augmented generation (RAG)**. It ingests an organisation's documents, classifies them against a vocabulary the organisation's own subject-matter experts control, indexes them for hybrid vector and keyword search, and answers questions in a chat interface, either its own or embedded in an existing intranet or customer portal.

What sets it apart is that **scope of visibility is part of the question**. Two people asking the same thing may legitimately get different answers, because they are entitled to different documents:

Dimension	Example values	How it is matched
Department	HR, Finance, Sales, IT, Legal, Support	The user needs any one of the document's departments
Region	Global › EMEA › UK, DE · AMER › US · APAC › IN	Hierarchical: a UK user also sees EMEA and Global documents
Clearance level	Public, Internal, Confidential, Restricted	The document's level must be at or below the user's
Role	Admin, taxonomy editor (SME), reviewer, contributor	Governs what a person can administer, classify or upload
Personal share	One document for one named person	Matched on the user's directory object ID

The attribute set is not hard-coded. It lives in a declarative access policy, so an organisation can add cost centre, legal entity, project or any other attribute without changing code.

Who it is for: organisations of any size whose knowledge is spread across departments and must not be shared indiscriminately, starting with mid-sized enterprises on Microsoft Azure and growing toward SMBs through a subscription offering.

3 High-level product overview

RAG-OS has two independent paths that meet at a shared search index. The **ingestion path** turns documents into classified, searchable passages; it runs on its own compute and scales from zero to many workers without affecting users. The **query path** turns a signed-in user's question into a cited answer drawn only from passages that user is entitled to see.





Figure 1. The two paths at a glance. Both use the same embedding profile, so a question is always compared with documents in the same vector space.

RAG-OS system architecture

Container Apps environment cae-ragos-dev - every app and job authenticates with the user-assigned managed identity id-ragos-dev (no keys, no connection strings)



Application Insights + Log Analytics (appi-ragos-dev - log-ragos-dev)

Every app and job exports OpenTelemetry traces, logs and metrics. One X-Correlation-ID follows a request from the browser through the API, onto the queue message and into the worker. Metrics: ragtokens - rag.stage.duration - rag.ingest.docs - KQL in docs/kql.md

Answering a question

- 1 The user signs in with Microsoft; MSAL acquires an access token for the RAG-OS API scope.
- 2 The page calls /api/* with that token; nginx adds X-Correlation-ID and proxies to the internal API.
- 3 rag-api validates issuer/algorithm/audience/expiry and maps the claims to the caller's attributes.
- 4 The question is embedded by rag-embed-query (tei only), or by Foundry text-embedding-3-small (aoai only).
- 5 Hybrid + semantic search runs with the access filter INSIDE the query (never a filter afterwards).
- 6 Claude or GPT answers from the numbered passages; citations and token usage come back with the answer.

Ingesting a document

- 7 Scheduler lists each due source and detects what changed (etag / content hash).
- 8 Changed documents are staged to Blob and enqueued as claim-check messages.
- 9 KEDA scales workers on queue length; the priority lane is drained first.
- 10 Chunks are embedded on the GPU pool (tei only); facets come from those vectors.
- 11 New chunks are upserted, the previous version is deleted, and the state is recorded.

Figure 2. System architecture: query path, ingestion path and shared Azure services.

3.1 Data ingestion path

Capability	Details	Status
Configurable data sources	Local folder, Azure Blob Storage, and direct upload through the UI or API (served on a priority lane). Sources are declared in configuration and built by a factory, so adding one is a plug-in.	AVAILABLE
Further connectors	SharePoint (with permission sync), Google Drive, SFTP/FTPS. Registered in the product, not yet implemented.	ROADMAP
File formats	PDF (one section per page), DOCX, TXT/Markdown/log, CSV/TSV, XLSX/XLSM, XLS, JSON, JSONL/NDJSON, JSONP, XML.	AVAILABLE
Scanned documents (OCR), PPTX, legacy DOC	Planned through Azure Content Understanding.	ROADMAP
Classification	Facets with owners, synonyms and hierarchies. Applied in order: folder path rules, a manifest, per-file sidecar metadata, then an automatic classifier. Low-confidence results go to an SME review queue.	AVAILABLE
Indexing	Azure AI Search: hybrid vector and keyword search, semantic ranker, int8 vector compression. Access attributes are stored on every passage.	AVAILABLE
Scale and resilience	Queue-driven (Azure Service Bus) and idempotent, with a per-document state machine, retries, a dead-letter queue and a status report per document. Designed for millions of documents.	AVAILABLE

How documents get classified

Cheapest layer first. Later layers override earlier ones for the same facet; an SME decision is never overwritten by a rule. Only facets still empty reach the classifier.




Figure 3. How documents are organised and classified, layer by layer. SMEs own the vocabulary; the system applies it.

3.2 User query path

1. **Sign in.** The user signs in with Microsoft Entra ID, in the RAG-OS chat UI or in a host portal that embeds it with a single script tag.
2. **Resolve scope.** The token's claims are mapped to the user's department, region, clearance and roles. Anything missing defaults to *deny*.
3. **Filter inside the search.** The access policy is compiled into a search filter and applied inside Azure AI Search, so only permitted passages are ever candidates.
4. **Retrieve and rank.** Hybrid vector and keyword retrieval is followed by semantic re-ranking.
5. **Answer or abstain.** The language model answers only from the retrieved passages and cites them as `[n]`. If there is nothing relevant, or the answer would be uncited, RAG-OS declines and says why.

What happens when someone asks a question

Every step carries the same X-Correlation-ID, so one identifier links the browser request, the search call and the model call in Application Insights.




Figure 4. A question travels from a signed-in browser to a cited answer.

3.3 Answer quality scenarios

Four behaviours separate a trustworthy assistant from a plausible one. Each is stated with its current status.

Version correctness

PARTIAL

When a document is re-ingested, its previous version is removed from the index, so stale copies of the *same* document are never quoted.

Choosing between *different* documents (for example `ProductPriceCatalog_2025` vs `ProductPriceCatalog_2026`) currently relies on an instruction to prefer the most recent effective date. The effective date is read from Markdown and text metadata today.

Roadmap: effective dates for every format, version-aware ranking, and an explicit "as of" filter in the query.

Page-level citation

AVAILABLE

Every answer carries numbered citations with the document, the **page number** (for PDFs) or section heading, a snippet and a relevance score.

When the same passage appears in several files, the citation lists the other locations too, instead of hiding them.

Example: "Parental leave is 26 weeks [1]." → [1] UK Parental Leave Handbook, page 4.

Abstention

AVAILABLE

RAG-OS returns no answer rather than a plausible one, with a machine-readable reason:

- `no_access`: nothing this user may see
- `no_relevant_context`: the corpus does not cover it
- `model_refusal`: the model declined
- Grounding guard: an answer without citations is discarded

Conflict handling

PARTIAL

When two permitted sources disagree, for example a global and a UK handbook, the answer surfaces the conflict and cites both, rather than choosing silently.

A user who can see only one of the two sources gets a clean answer from that source, because the other was never retrieved.

Roadmap: a structured `conflict` flag in the API response, so portals and workflows can act on it.

4 Models

RAG-OS treats models as configuration, not code. Two model roles are configured independently.


Role	Options	Status
Embeddings turn text into vectors	• Default: Qwen3-Embedding-0.6B (Apache-2.0), self-hosted on Hugging Face Text Embeddings Inference in your own containers. Multilingual, 1024 or 512 dimensions, no per-token cost, and text never leaves your environment. • Alternative: Azure OpenAI <code>text-embedding-3-small</code>.	AVAILABLE
Answering compose the cited answer	• Claude models on Microsoft Foundry, or Azure OpenAI GPT models. Both use keyless Entra authentication. • A separate, cheaper "utility" model for background tasks, and an optional fallback provider.	AVAILABLE
On-premises / confidential compute	Adapters for models hosted on-premises or on confidential-compute GPUs. The provider registry is already pluggable; the adapters are not yet built.	ROADMAP

The embedding-profile guarantee. Exactly one embedding profile is active at a time. It pins the model, revision, dimensions and query/document prefixes, and both ingestion and querying read the same profile. A mismatch between how documents and questions are embedded, a common and silent cause of poor RAG quality, cannot happen. Changing the profile builds a new index alongside the old one, which keeps serving until the new one is ready.

5 Secure data querying

Users are governed by **Microsoft Entra ID**: app roles decide what a person can *do*, and directory attributes decide what a person can *see*. Documents carry matching facets (tags) in the index. A declarative policy joins the two.

Clearance levels




Levels are configurable in the access policy, so an organisation that needs an extra tier, such as "Private", can add one.

Pre-filtering, not post-filtering

Many assistants retrieve broadly and then hide what the user should not see. RAG-OS never retrieves it. The policy (default deny; department *and* region *and* clearance, *or* a personal share) is compiled into a filter that the search engine applies **inside** every query. Restricted text never reaches the model, the logs or the answer.

How a person's attributes become a search filter

Worked example: Priya — HR, United Kingdom, clearance 1, employee id E1001, signed in with Entra ID. Adding an attribute is a YAML edit; none of this is hard-coded. Three names for one idea: (1) the CLAIM Entra sends about the person, (2) the ATTRIBUTE the policy calls it, (3) the `acl_*` INDEX FIELD on the document. Nothing named `acl_*` is ever a claim.



A property test generates random principals and documents and checks that this OData filter and the Python predicate used by the local index always agree — so what you test offline is what Azure enforces.

Figure 5. How a person's token claims become the filter applied inside every search.

Control	What it does	Status
Entra ID sign-in and app roles	Admin, taxonomy editor (SME), reviewer and contributor roles	AVAILABLE
Attribute administration	Admins set a person's department, region and clearance from the admin console, written to Entra through Microsoft Graph	AVAILABLE
"Explain access"	Shows exactly why a given user can or cannot see a given document	AVAILABLE
Prompt-injection hygiene	Document content is treated as data, never as instructions	AVAILABLE
Secrets	Azure Key Vault and managed identities; no secrets in code, images or files	AVAILABLE
Source permission sync	Inherit SharePoint ACLs in addition to the attribute policy	ROADMAP
Microsoft Purview labels	Map sensitivity labels to clearance levels automatically	ROADMAP
Private networking	Private endpoints and VNet integration	ROADMAP

6 Data compliance

For regulated organisations the first question is not "how good are the answers" but "where does our data go". RAG-OS is designed to keep the answer to that question short.

Your tenant, your subscription

AVAILABLE

RAG-OS deploys into the customer's own Azure subscription. Documents, index, database, logs and keys stay under the customer's control and policies.

Embeddings never leave

AVAILABLE

With the default self-hosted embedding model, document text is vectorised inside the deployment. No document content is sent to a third-party embedding API.

Minimal exposure to the LLM

AVAILABLE

Only the handful of passages the user is already entitled to see are sent to the answering model, through Microsoft Foundry or Azure OpenAI in the customer's chosen region.

Compliance-perimeter models

ROADMAP

Answering models running on-premises or on Azure confidential-compute infrastructure, so that even the answering step stays inside a defined perimeter.

7 Observability and operations

Tokens per answer

AVAILABLE

Every answer reports input, output, cache and embedding tokens, split by purpose, and the same figures are emitted as metrics by provider and model.

End-to-end tracing

AVAILABLE

A correlation ID follows a request from the browser through the API, and follows a document from upload through queue, worker and index. OpenTelemetry exports to Azure Application Insights.

Ready-made queries

AVAILABLE

A library of KQL queries for traces, token dashboards, failures and alerts ships with the product.

Admin console

AVAILABLE

Dashboard, sources and "sync now", ingestion runs with progress and ETA, per-document status, bulk retry, CSV export, SME review queue, dead-letter viewer, ingestion throttling and configuration editing.

8 How RAG-OS is different

Compared with Azure knowledge bases, NotebookLM and similar RAG products, RAG-OS competes on governance and control rather than on connector count.

Differentiator	Why it matters	Status
Attribute-based access enforced inside the search	Department, region, clearance and personal shares are matched before retrieval. Nothing is retrieved and then hidden.	AVAILABLE
Clearance levels and region hierarchy	Models how regulated organisations actually classify information, beyond simple per-file sharing.	AVAILABLE
Explainable access	Admins can answer "why can this person see this?" in seconds, which matters for audits.	AVAILABLE
Configurable models	Self-hosted embeddings, Claude or GPT for answering. No lock-in to one model vendor.	AVAILABLE
Honest abstention	Reason-coded refusals and a grounding guard, instead of plausible guesses.	AVAILABLE
Cost transparency	Token usage on every answer, and zero-token-cost embeddings.	AVAILABLE
SME-owned taxonomy	Subject-matter experts own the vocabulary and review queue; IT owns the platform.	AVAILABLE
Embeddable anywhere	One script tag puts governed chat inside an intranet or customer portal, scoped by the host's identity token.	AVAILABLE
SaaS model to scale	A multi-tenant offering for SMBs alongside dedicated deployments for enterprises.	ROADMAP
Custom evaluation pipeline	Combines RAGAS, DeepEval and Azure AI Evaluation metrics (faithfulness, answer relevance, context precision and recall) and gates releases on them.	ROADMAP

Suggested additions to stand out

- **Answer confidence score** **ROADMAP**: a per-answer faithfulness score from a hallucination-detection model, shown to users and usable as a threshold.
- **Structured version and conflict signals** **ROADMAP**: "as of" dates and conflict flags in the API, not just in the answer text.
- **Governed retrieval as an MCP server** **ROADMAP**: let any agent, such as Copilot or a Foundry agent, use RAG-OS retrieval while keeping its access policy.
- **Domain packs** **ROADMAP**: pre-built taxonomy, ontology and evaluation sets per industry, such as mortgage servicing, insurance claims or HR, to shorten time-to-value.
- **Answer audit trail** **ROADMAP**: an exportable record of who asked what and which documents informed the answer, for regulators.

9 Similar products in market

The enterprise knowledge-assistant market is crowded and moving quickly. The products below are the ones a buyer is most likely to compare RAG-OS with. Positioning reflects public information as of October 2026.

Product	What it is	Where RAG-OS differs
Microsoft 365 Copilot / Copilot Studio	Assistant across Microsoft 365 content, with Purview sensitivity labels shown on citations and inherited by agents.	Works with content outside Microsoft 365, applies a custom attribute policy (region, clearance), and can serve external portal users.
Foundry IQ (Azure AI Search knowledge bases)	A knowledge layer for agents: agentic retrieval over Azure AI Search with document-level permissions and an MCP server.	A building block, where RAG-OS is a finished product with ingestion operations, an SME taxonomy, an admin console and a chat UI. RAG-OS can adopt Foundry IQ underneath (roadmap).
Glean	A work-AI platform with a permissions-aware knowledge graph over 100+ SaaS apps. Strong in large enterprises.	Glean mirrors source ACLs; RAG-OS adds clearance and region policy independent of the source, runs in the customer's own Azure subscription, and costs less to run.
Google Gemini Enterprise / Gemini Notebook Enterprise (formerly NotebookLM Enterprise)	Google Cloud's agent platform with permission-aware connectors; notebooks shared per notebook.	Azure-native; organisation-wide attribute policy instead of per-notebook sharing.
Amazon Q Business / Amazon Quick Suite	AWS's permission-aware enterprise assistant, now evolving into the agentic Quick Suite.	For Azure- and Microsoft-centric organisations; configurable models, including self-hosted embeddings.
Atlassian Rovo	Search, chat and agents over Jira, Confluence and 50+ connected apps; acts strictly within each user's permissions.	Not tied to a collaboration suite; designed for policy and handbook corpora with clearance levels.

Hebbia (Matrix)	Agentic document analysis for finance and legal, used by leading asset managers.	Horizontal and affordable; reaches verticals through SME taxonomy and ontology rather than a premium vertical product.
Vectara	RAG-as-a-service platform known for hallucination detection (the HHEM model) and runtime correction.	A complete governed product in the customer's tenant. A hallucination score is a candidate addition to RAG-OS's evaluation roadmap.
Onyx (formerly Danswer)	Open-source enterprise search and chat with 40+ connectors and source-permission sync; self-hostable.	Azure-native operations at scale, attribute and clearance policy, and SME classification with a review queue.

Positioning in one line: the market leaders compete on how many systems they connect to. RAG-OS competes on how precisely it controls who sees what, and on proving where every answer came from, inside infrastructure the customer already trusts.

10 Potential use cases

The core product is a governed question-answering service. Because the knowledge base is indexed, classified and permission-tagged, the same foundation supports a family of higher-value applications.

10.1 Knowledge-base assistant for SMBs (SaaS subscription) CORE AVAILABLE

SAAS ROADMAP

Most SMBs keep their standard operating procedures across HR, Sales, IT, Legal, Finance and Training in scattered folders. RAG-OS makes them queryable, with reliable, cited answers and department-level privacy, as a subscription that needs no AI expertise to run.

- **Packaging:** tiers by seats and document volume; usage-based add-ons, metered by the token figures the product already records.
- **Onboarding:** connect a folder or storage account, apply a starter taxonomy, and invite users.

10.2 Questionnaire generation for multiple verticals ROADMAP

A searchable, classified knowledge base can be turned into assessments as well as answers. Every generated question is grounded in, and cites, a source passage.

Question types

- Objective (multiple choice)
- Fill in the blanks
- Match the columns
- Descriptive / open questions

Verticals

- Education: quizzes and exams from syllabus material
- Hiring: online and offline interview sets from role profiles
- Compliance and training certification
- Fun quizzes and engagement

Evaluation

Objective questions are marked automatically. Descriptive answers are scored against the cited source passages, and the evaluator's confidence decides how marks are allocated; low-confidence results go to a human reviewer.

10.3 Form filling ROADMAP

Use the indexed knowledge base to pre-fill structured forms, with a citation behind every value so a person can verify it before submitting.

- Legal forms and standard agreements
- Tax forms
- Application forms: licences, permits, grants, admissions, visas
- Custom templated reports for compliance and audits

10.4 Personal assistants ROADMAP

Individuals keep educational, financial, employment, property and household documents in one private, searchable space ("When does my home insurance renew?", "What was my gross salary in 2025?"). Combined with form filling, the personal knowledge base becomes the data source for everyday paperwork.

10.5 Niche verticals: finance, insurance, healthcare ROADMAP

The highest-value opportunities are narrow, rule-heavy processes where checking documents against a specific *version* of a specific guideline consumes skilled staff time. The same pattern applies to insurance claims, KYC and onboarding files, clinical prior-authorisation packets and grant applications. Subject-matter experts are the ones who spot these niches; RAG-OS gives them the tools to encode their judgement.

Worked example: mortgage loss-mitigation document chase. A homeowner falls behind on payments and applies for hardship assistance. The servicer must decide whether the submitted package is complete against the relevant investor guideline, then state precisely what is missing. The judgement is *completeness against a specific version of a specific guide*.

How RAG-OS addresses it:

1. **Versioned guideline corpus.** Each investor guide is indexed with its effective dates, so the check runs against the version in force for the application.
2. **Taxonomy.** SMEs define the document types: hardship statement, proof of income, bank statements, tax records, property and occupancy evidence.
3. **Ontology.** SMEs define the relationships: *guideline version* → *requires* → *document type*, with conditions (for example, a self-employed borrower requires profit-and-loss evidence) and validity rules such as how recent a statement must be.
4. **Completeness agent.** The agent classifies the submitted package, evaluates it against the ontology for that guideline version, and produces a missing-items list where every requirement cites the guideline page it comes from.

11 Future scope and potential

The roadmap moves in three steps: complete the enterprise product, then open the platform, then build workflows on top of it.

Next

Complete the enterprise product

- SharePoint connector with permission sync, then Google Drive and SFTP
- OCR for scanned documents (Azure Content Understanding)
- Evaluation pipeline (RAGAS, DeepEval, Azure AI Evaluation) as CI/CD release gates
- Version-aware ranking, and structured conflict and "as of" signals
- UI editors for taxonomy and access policy
- Streaming answers; Bicep templates; private endpoints and VNet

Later

Open the platform

- Multi-tenant SaaS control plane: tenant isolation, per-tenant configuration, metering and billing on recorded token usage
- On-premises and confidential-compute model adapters
- Publish as a Foundry agent for Teams and Copilot Studio; expose governed retrieval as an MCP server
- Foundry IQ knowledge bases as an alternative retrieval engine
- Purview sensitivity-label mapping; semantic cache and an API-gateway token budget

Horizon

Workflows on governed knowledge

- Questionnaire generation and grading
- Citation-backed form filling
- Completeness-checking agents for regulated verticals
- Ontology and knowledge-graph reasoning across documents
- A marketplace of domain packs (taxonomy, ontology and evaluation sets)
- Personal knowledge assistant edition

Market potential

Three revenue streams build on one platform:

- **Dedicated enterprise deployments** for regulated organisations that need the product inside their own subscription.
- **SaaS subscriptions** for SMBs that need a trustworthy SOP assistant without an AI team.
- **Vertical solutions** (domain packs and workflow agents) priced on the effort they save, such as hours of document chasing per mortgage file.

12 Appendix

A. Technology stack

Component	Choice	Purpose
API and workers	Python 3.13, FastAPI, pydantic	One codebase for the API, ingestion workers and CLI; OpenAPI-documented
Embeddings	Hugging Face TEI + Qwen3-Embedding-0.6B	Self-hosted, multilingual vectors with no token cost
Answering LLM	Claude on Microsoft Foundry, or Azure OpenAI	Grounded, cited answer generation
Search	Azure AI Search	Hybrid vector and keyword retrieval with in-query security filters
Queue	Azure Service Bus	Reliable, idempotent ingestion with priority lanes and a dead-letter queue
State	Azure Database for PostgreSQL	Per-document status, runs and reporting at millions of rows
Files and configuration	Azure Blob Storage	Raw documents, versioned configuration, exports
Hosting	Azure Container Apps (incl. serverless GPU)	Separate, independently scaling query and ingestion compute
Identity and secrets	Microsoft Entra ID, Key Vault, managed identity	Sign-in, roles, attributes; no stored secrets
Observability	OpenTelemetry → Application Insights	Traces, logs and token metrics
Chat UI	TypeScript, MSAL, nginx	Chat, embeddable loader and admin console

B. Glossary

RAG	Retrieval-augmented generation: answering with a language model from passages retrieved at question time, rather than from the model's memory.
ABAC / RBAC	Attribute-based and role-based access control. RAG-OS uses roles for what a person can do and attributes for what a person can see.
Facet	A controlled tag on a document, such as department, region or clearance, used for filtering and classification.
Taxonomy	The SME-owned vocabulary of facets, values, synonyms and hierarchies.
Ontology	The relationships between concepts in a domain, for example "guideline version requires document type".
Embedding profile	The pinned model, revision, dimensions and prefixes used for both ingestion and queries.
Pre-filter	An access filter applied inside the search query, so that forbidden content is never retrieved.
Abstention	Declining to answer, with a reason, when the permitted evidence does not support an answer.
MCP	Model Context Protocol: an open standard for connecting AI agents to tools and data.

RAG-OS Product Overview · v1.0 · 2 October 2026. Statuses reflect the product as built at this date. Market information is drawn from public sources and may change.